

Machine learning encoding & decoding; linear & non-linear models

Mads Jensen, PhD

✉ mads@cas.au.dk



AARHUS UNIVERSITY



IMC
INTERACTING MINDS CENTRE



Contents

1. Machine learning
2. Encoding
 - example: rapid tuning shifts in human auditory cortex ...
3. Decoding
 - evaluating algorithms
 - linear vs non-linear models
4. Time generalisation
 - example: brain mechanisms underlying the brief maintenance ...
5. General paper feedback
 - abstract
 - parts of a paper
 - arguments
 - material condition

Machine learning

Machine learning: what is it?

- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning

Machine learning: what is it?

- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning
- learning from data

Machine learning: what is it?

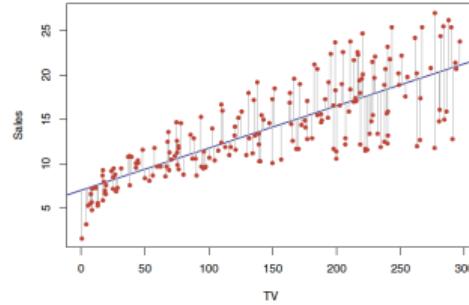
- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning
- learning from data
- $Y = f(X) + \epsilon$

Machine learning: what is it?

- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning
- learning from data
- $Y = f(X) + \epsilon$
- out-of-sample generalization

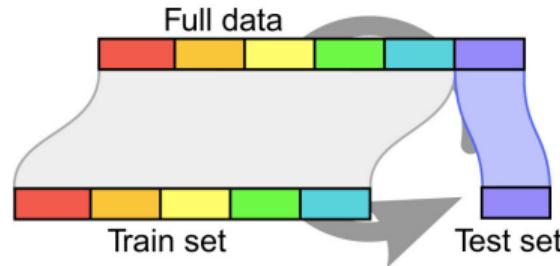
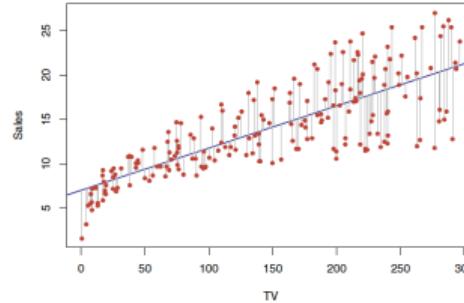
Machine learning: what is it?

- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning
- learning from data
- $Y = f(X) + \epsilon$
- out-of-sample generalization



Machine learning: what is it?

- multivariate pattern analysis,
multivoxel pattern analysis,
predictive modelling,
statistical learning
- learning from data
- $Y = f(X) + \epsilon$
- out-of-sample generalization



Top figure from James et al. (2013)

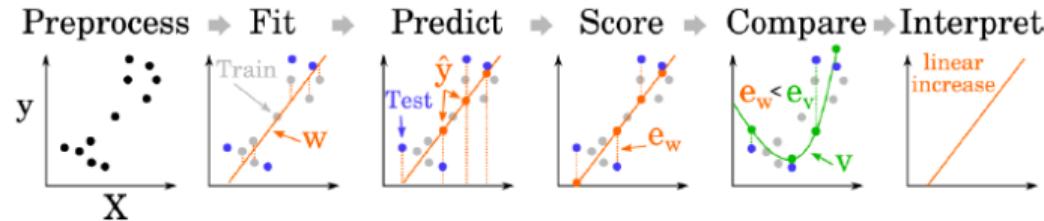
bottom figure from Varoquaux et al. (2017)

Machine learning: pipeline

- create features
- make cross-validation scheme
- fit model
- interpret model

Machine learning: pipeline

- create features
- make cross-validation scheme
- fit model
- interpret model



(Figure from King et al., 2018)

Features

What are the question to be investigated?

Features

What are the question to be investigated?

- whole brain (all sensors) vs ROIs

Features

What are the question to be investigated?

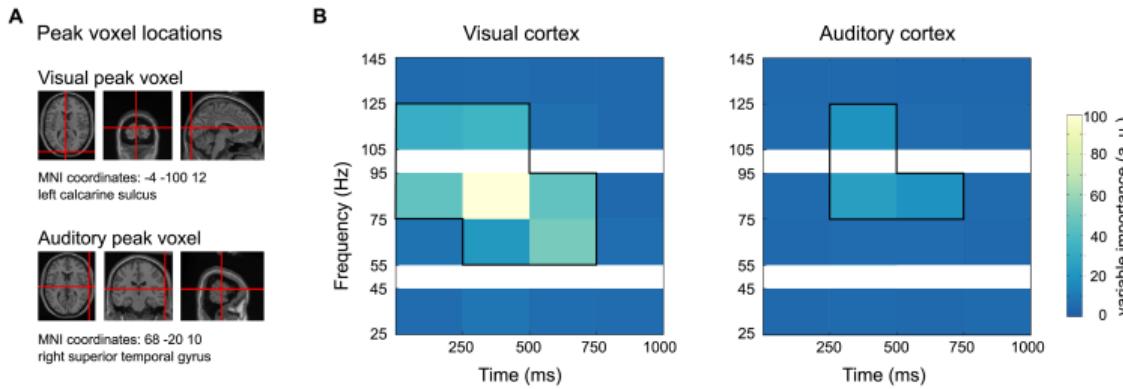
- whole brain (all sensors) vs ROIs
- raw signals vs transformed signals

Features

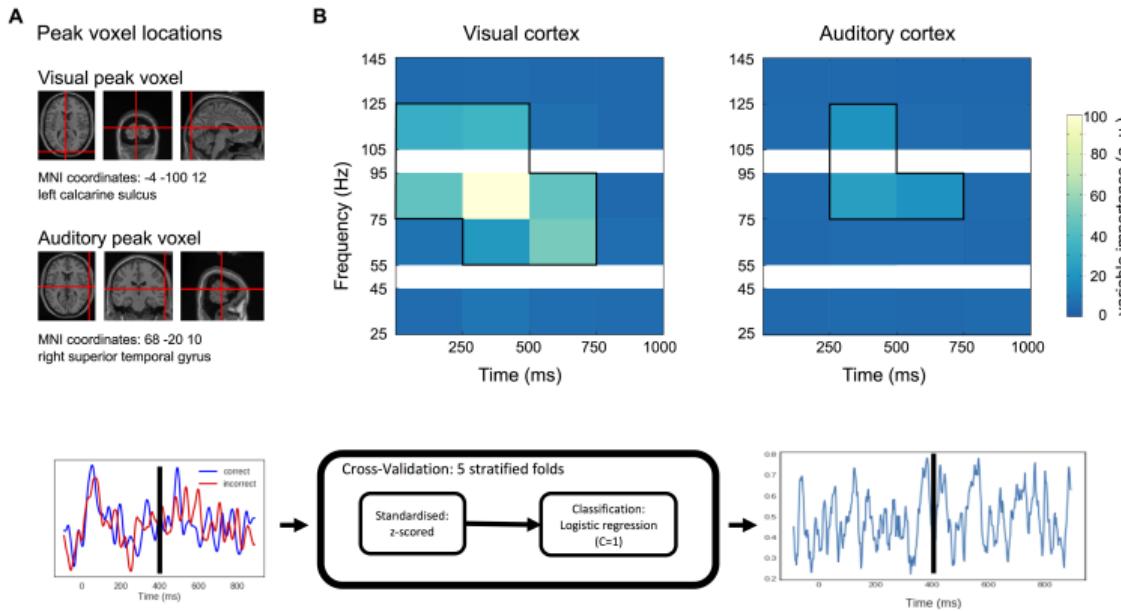
What are the question to be investigated?

- whole brain (all sensors) vs ROIs
- raw signals vs transformed signals
- all times together vs time bins vs across time

Features



Features



Top figure from Westner et al. (2018)

bottom figure mine.

Encoding

Encoding

Encoding is the attempt to model brain activity based on stimuli.

Encoding

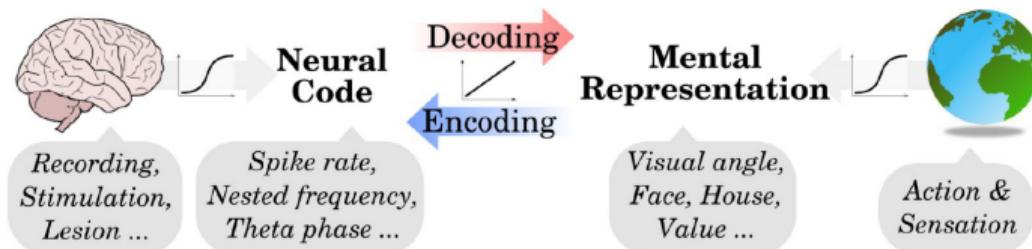
Encoding is the attempt to model brain activity based on stimuli.

- That is, given a specific stimulus, e.g. a sound, can we then predict the how the brain reacts.
- Often used algorithm is a linear regression
- Encoding is a world to brain relation.

Encoding

Encoding is the attempt to model brain activity based on stimuli.

- That is, given a specific stimulus, e.g. a sound, can we then predict the how the brain reacts.
- Often used algorithm is a linear regression
- Encoding is a world to brain relation.



(Figure from King et al., 2018)

Encoding

Linear regression problem

$$\text{activity}(t) = \sum_i^{N_{\text{features}}} \text{features}_i(t) * \text{weight}_i + \text{error}(t)$$

Where the neural activity at time t is modelled as a weighted sum of N stimulus features.
(Holdgraf et al., 2017, p. 5)¹

¹ Adapted from Hastie et al. (2009)

Encoding

Linear regression problem

$$\text{activity}(t) = \sum_i^{N_{\text{features}}} \text{features}_i(t) * \text{weight}_i + \text{error}(t)$$

Where the neural activity at time t is modelled as a weighted sum of N stimulus features.
(Holdgraf et al., 2017, p. 5)¹

Time lagged version

$$\text{activity}(t) = \sum_j^{N_{\text{lags}}} \sum_i^{N_{\text{features}}} \text{features}_i(t - j) * \text{weight}_{i,j} + \text{error}(t)$$

¹ Adapted from Hastie et al. (2009)

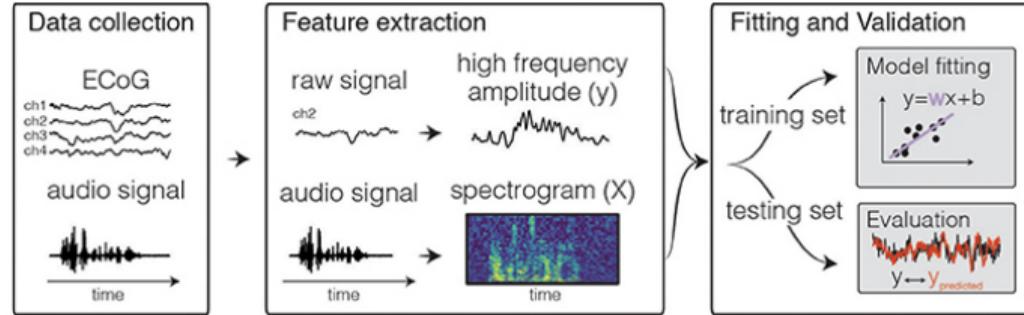
Encoding

Linear algebra terms:

$$\text{activity} = S\mathbf{w} + \epsilon$$

"In this case S is the stimulus matrix where each row corresponds to a timepoint of the response, and the columns are the feature values at that timepoint and time-lag (there are $N_{lags} * N_{features}$ columns). \mathbf{w} is a vector of model weights (one for each feature * time lag), and ϵ is a vector of random noise at each timepoint (most often to be Gaussian for continuous signals or Poisson for discrete signals)." (Holdgraf et al., 2017, p. 5)

Encoding



(Figure from Holdgraf et al., 2017)

ARTICLE

Received 27 May 2016 | Accepted 20 Oct 2016 | Published 20 Dec 2016

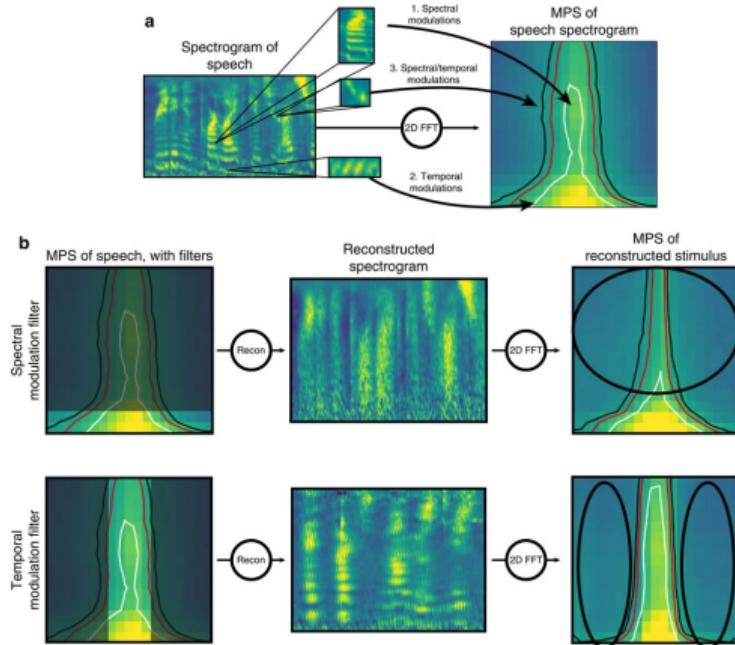
DOI: 10.1038/ncomms13654

OPEN

Rapid tuning shifts in human auditory cortex enhance speech intelligibility

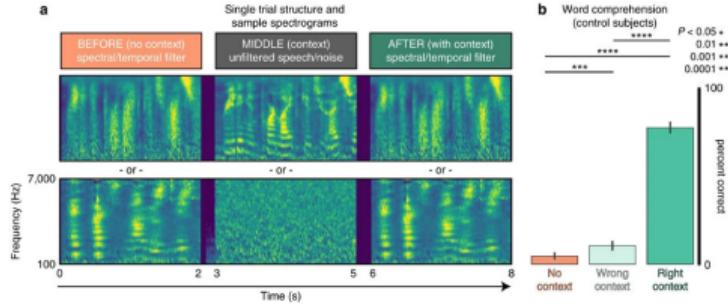
Christopher R. Holdgraf¹, Wendy de Heer², Brian Pasley¹, Jochem Rieger¹, Nathan Crone³, Jack J. Lin⁴, Robert T. Knight^{1,2,3} & Frédéric E. Theunissen^{1,2}

Rapid tuning shifts in human auditory cortex . . .



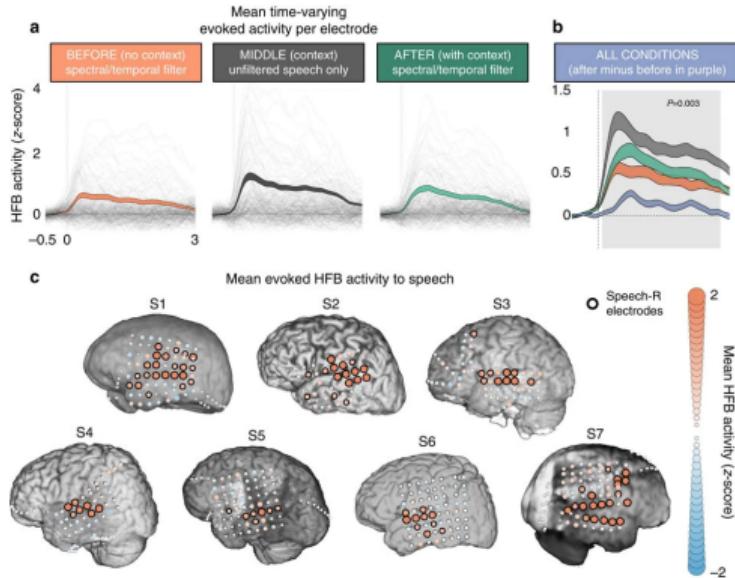
(Figure from Holdgraf et al., 2016)

Rapid tuning shifts in human auditory cortex . . .



(Figure from Holdgraf et al., 2016)

Rapid tuning shifts in human auditory cortex . . .



(Figure from Holdgraf et al., 2016)

Decoding

Decoding

Decoding is the attempt to predict stimuli based on brain activity.
Decoding is also called *classification*.

Decoding

Decoding is the attempt to predict stimuli based on brain activity.

Decoding is also called *classification*.

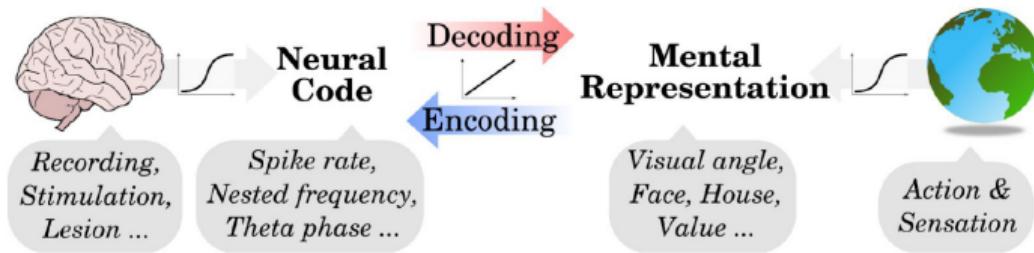
- That is, given a specific pattern of brain activity can we predict which stimuli the participant perceived, e.g. based on brain activity can we predict which sound the participant heard.
- Often used algorithm is a logistic regression or similar linear model.
- Decoding is a brain to world relation.

Decoding

Decoding is the attempt to predict stimuli based on brain activity.

Decoding is also called *classification*.

- That is, given a specific pattern of brain activity can we predict which stimuli the participant perceived, e.g. based on brain activity can we predict which sound the participant heard.
 - Often used algorithm is a logistic regression or similar linear model.
 - Decoding is a brain to world relation.

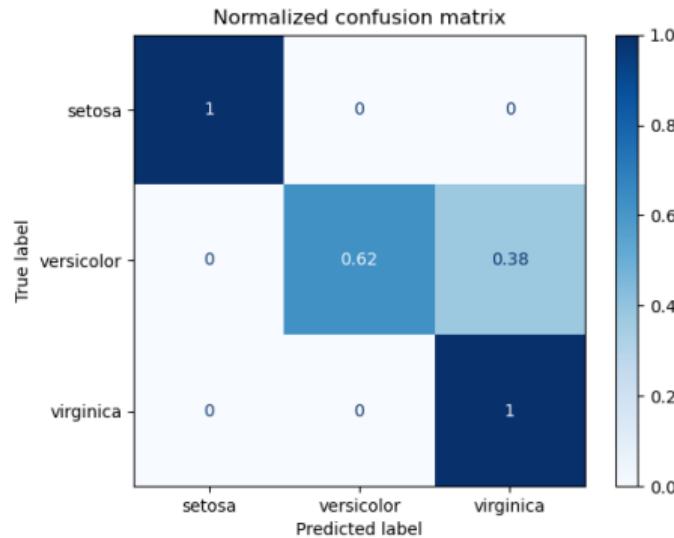
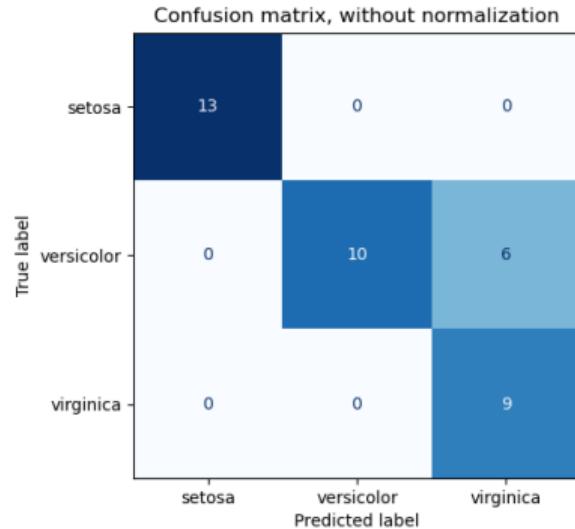


Decoding

$$\text{feature}(t) = \sum_j^{N_{\text{lags}}} \sum_i^{N_{\text{channels}}} \text{activity}_i(t+j) * \text{weight}_{i,j} + \text{error}(t)$$
$$s = Xw + \epsilon$$

“where s is a vector of stimulus feature values recorded over time, and X is the channel activity matrix where each row is a timepoint and each column is a neural feature (with time-lags being treated as a separate column each). w is a vector of model weights (one for each neural feature * time lag), and ϵ is a vector of random noise at each timepoint (often assumed to be Gaussian noise).” (Holdgraf et al., 2017, p. 5)

Confusion matrix



(Figure from https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html)

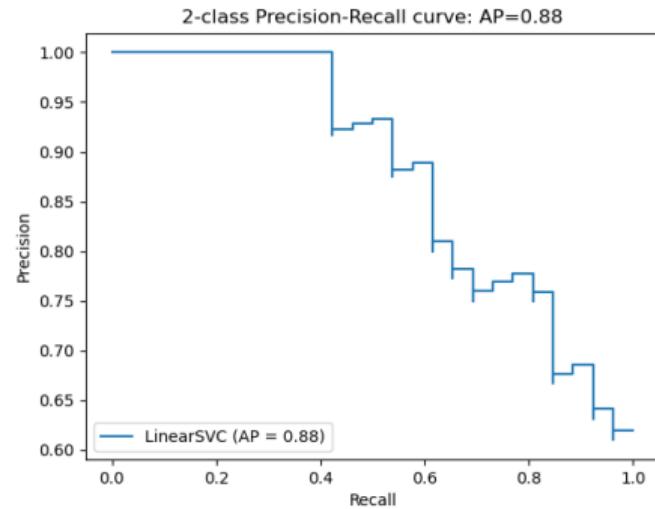
Precision-Recall

Precision (P) is defined as the number of true positives (T_p) over the number of true positives plus the number of false positives (F_p).

$$P = \frac{T_p}{T_p + F_p}$$

Recall (R) is defined as the number of true positives (T_p) over the number of true positives plus the number of false negatives (F_n).

$$R = \frac{T_p}{T_p + F_n}$$



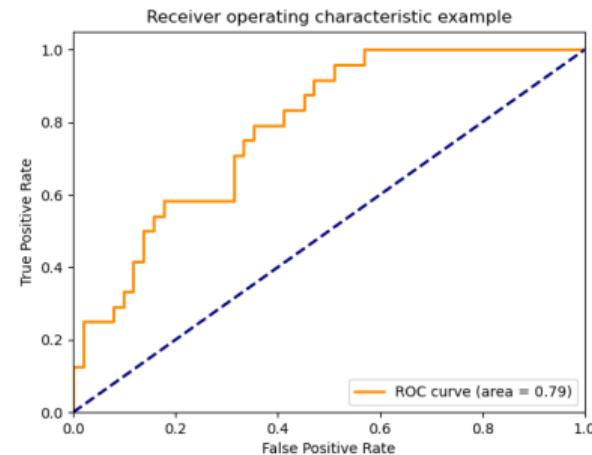
(Figure from https://scikit-learn.org/stable/auto_examples/model_selection/plot_precision_recall.html)

Receiver Operating Characteristic (ROC)

ROC are *true positive rate* (TPR; recall) plotted against the *false positive rate* (FPR)

$$TPR = \frac{T_p}{T_p + F_n}$$

$$FPR = \frac{F_p}{F_p + T_n}$$



(Figure from https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.html)

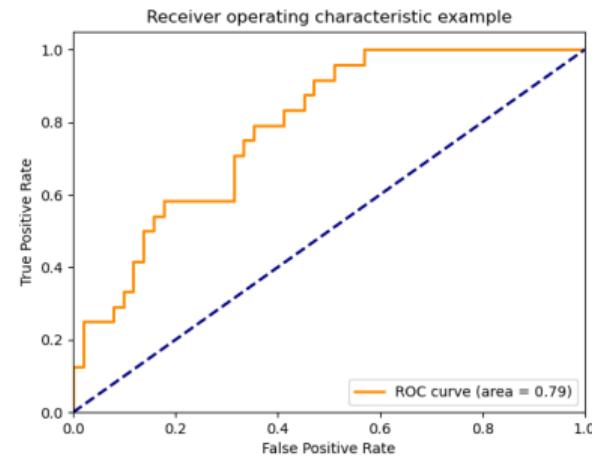
Receiver Operating Characteristic (ROC)

ROC are *true positive rate* (TPR; recall) plotted against the *false positive rate* (FPR)

$$TPR = \frac{T_p}{T_p + F_n}$$

$$FPR = \frac{F_p}{F_p + T_n}$$

"The Area Under the Curve (AUC) is simply the total amount of area under the ROC curve, and is often reported as a summary statistic of the ROC curve.



(Figure from https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.html)

Receiver Operating Characteristic (ROC)

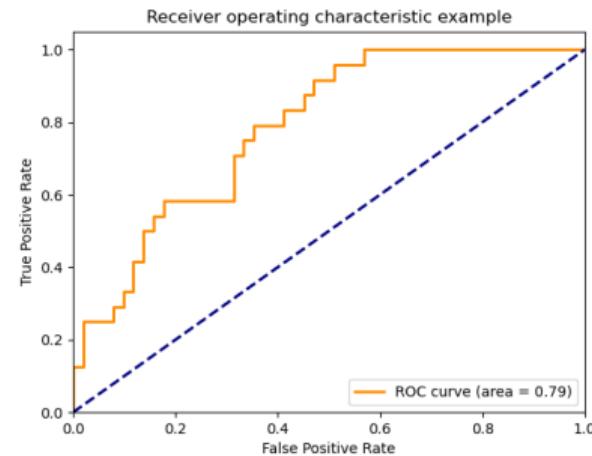
ROC are *true positive rate* (TPR; recall) plotted against the *false positive rate* (FPR)

$$TPR = \frac{T_p}{T_p + F_n}$$

$$FPR = \frac{F_p}{F_p + T_n}$$

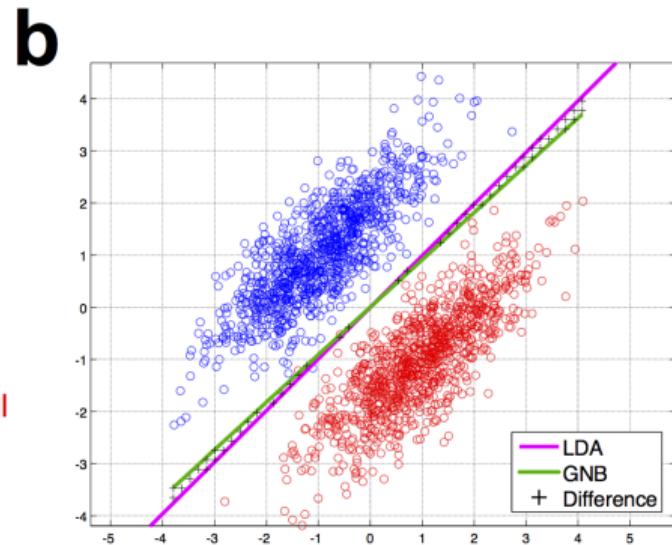
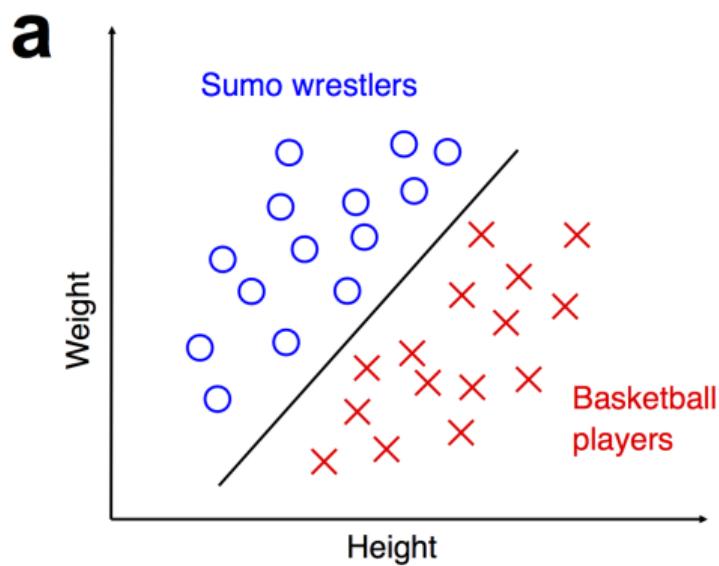
"The Area Under the Curve (AUC) is simply the total amount of area under the ROC curve, and is often reported as a summary statistic of the ROC curve.

If the classifier is performing at chance, then the AUC will be 0.5, and if it correctly labels all datapoints for all decision thresholds, then the AUC will be 1 " (Holdgraf et al., 2017, p. 15)



(Figure from https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.html)

Decision boundaries



(Figure from Raizada & Lee, 2013)

Regularisation

Objective:

$\max P(\mathbf{w} | X, y)$,
such that: $x_i^T \mathbf{w} = y_i$

Method:

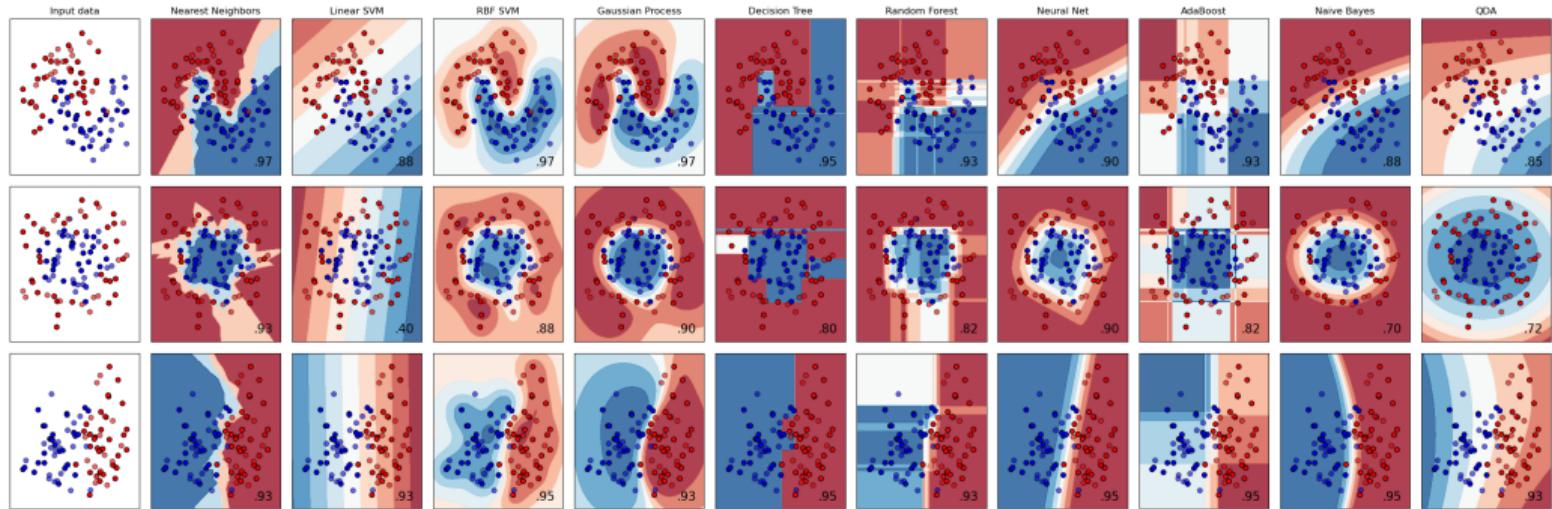
$P(\mathbf{w} | X, y) \propto P(X, y | \mathbf{w}) P(\mathbf{w})$
 $\min_{\mathbf{w}} \text{Loss}(y, X^T \mathbf{w}) + \text{Regularization}(\mathbf{w})$

Examples:

Name	Loss	Model	Reg	Optimization
OLS:	l_2		None	$\min_{\mathbf{w}} \sum_i (y_i - \hat{y}_i)^2$
Ridge:	l_2	$x_i^T \mathbf{w} = \hat{y}_i$	l_2	$\min_{\mathbf{w}} \sum_i (y_i - \hat{y}_i)^2 + \lambda \ \mathbf{w}\ ^2$
Logistic:	\log		l_2	$\min_{\mathbf{w}} \sum_i \log(1 + \exp(-y_i \hat{y}_i)) + \lambda \ \mathbf{w}\ ^2$
SVM:	<i>hinge</i>		l_2	$\min_{\mathbf{w}} \sum_i \max(0, 1 - y_i \hat{y}_i) + \lambda \ \mathbf{w}\ ^2$

(Figure from King et al., 2018)

Linear vs non-linear models



(Figure from https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html)

Linear vs non-linear models

Examples of non-linear classifiers:

- Random forest (Breiman, 2001)
- XGBoost (Chen & Guestrin, 2016)

Linear vs non-linear models

Examples of non-linear classifiers:

- Random forest (Breiman, 2001)
- XGBoost (Chen & Guestrin, 2016)

Pros of non-linear classifiers

- can estimate non-linear decision boundaries

Linear vs non-linear models

Examples of non-linear classifiers:

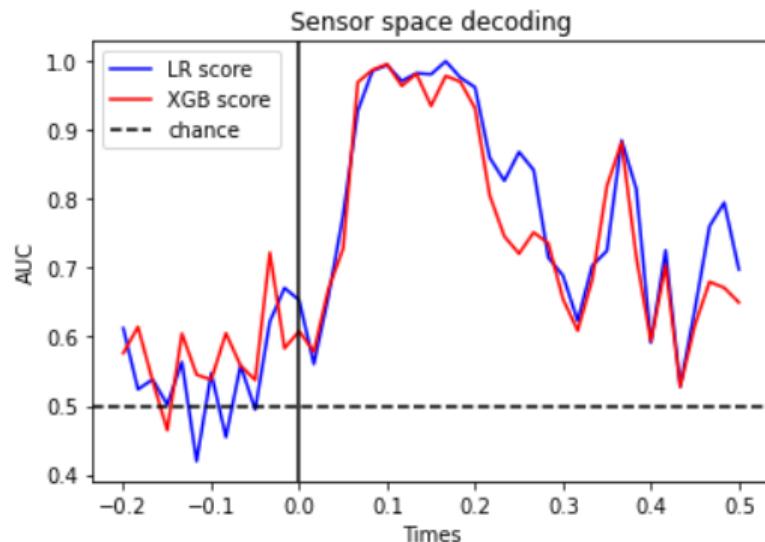
- Random forest (Breiman, 2001)
- XGBoost (Chen & Guestrin, 2016)

Pros of non-linear classifiers

- can estimate non-linear decision boundaries

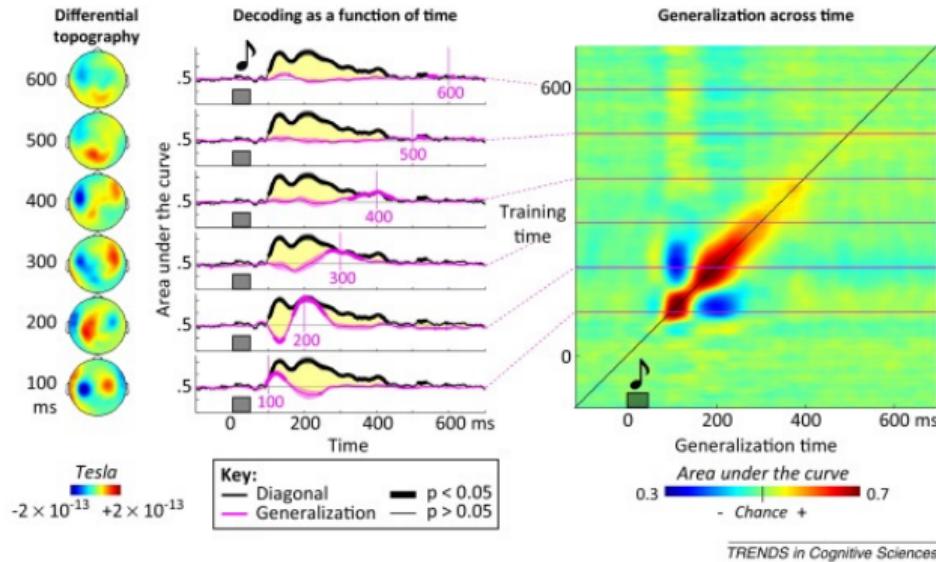
Cons non-linear classifiers

- require more data than linear classifiers
- require tuning of hyper-parameters
- often does not perform better
- hard to interpret



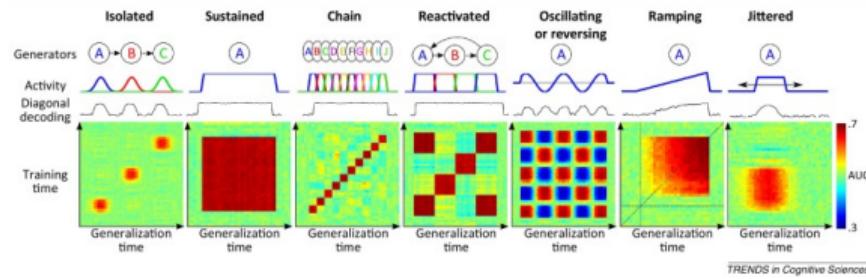
Time generalisation

Time generalisation

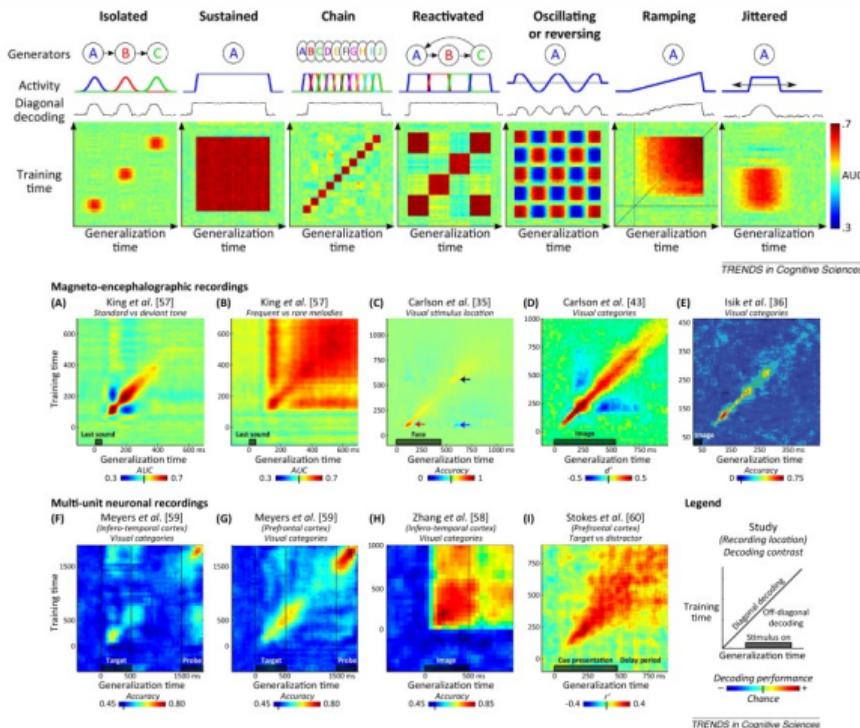


(Figure from King & Dehaene, 2014)

Time generalisation



Time generalisation



(Figure from King & Dehaene, 2014)

Mads Jensen (RFR, IMC, & CFIN)

machine learning

Brain Mechanisms Underlying the Brief Maintenance of Seen and Unseen Sensory Information

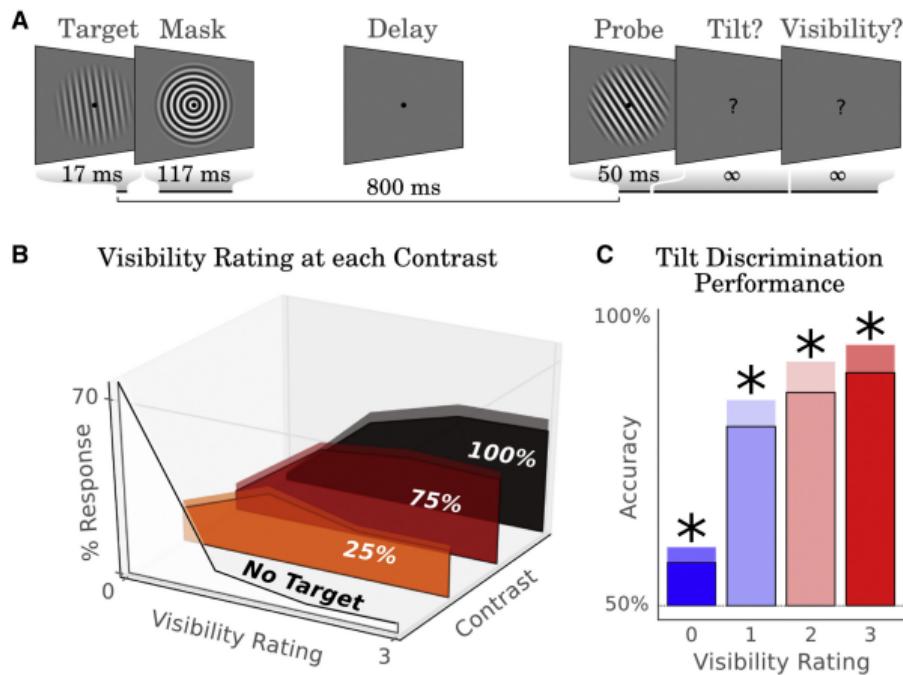
Jean-Rémi King,^{1,2,6,7,*} Niccolò Pescetelli,^{2,4} and Stanislas Dehaene^{4,5}

¹Department of Psychology, New York University, New York, NY 10003, USA

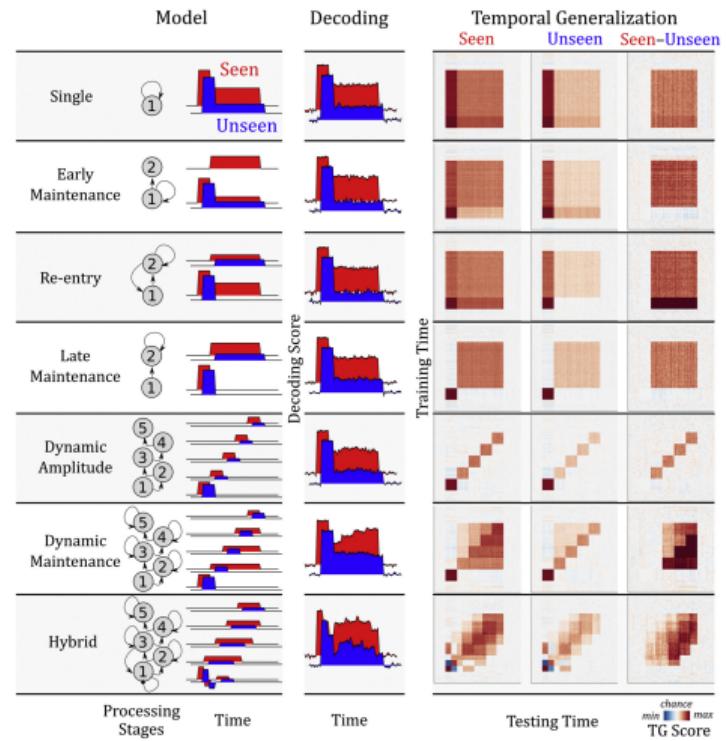
²Neuroscience Department, Frankfurt Institute for Advanced Studies, 60438 Frankfurt, Germany

³Department of Experimental Psychology, University of Oxford, OX1 3UD Oxford, UK

⁴Cognitive Neuroimaging Unit, CEA DSM/I2BM, INSERM, Université Paris-Sud, Université Paris-Saclay, NeuroSpin Center, 91191 Gif/Yvette, France



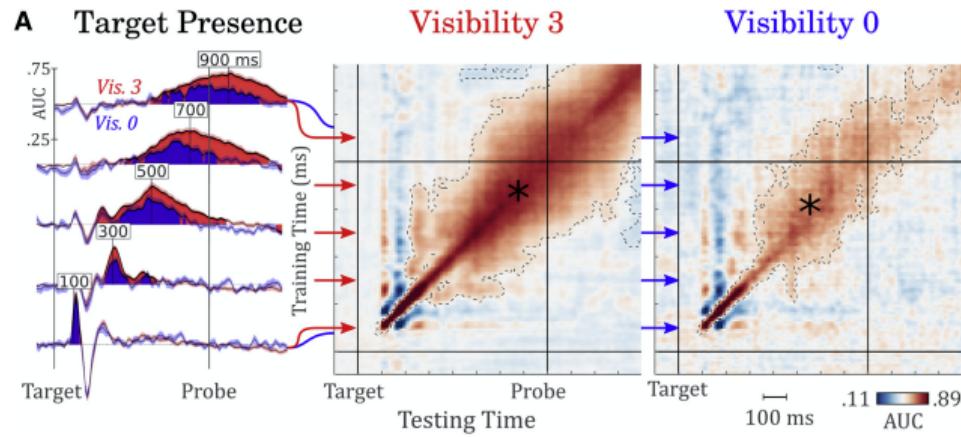
(Figure from King et al., 2016)



(Figure from King et al., 2016)

Mads Jensen (RFR, IMC, & CFIN)

machine learning



(Figure from King et al., 2016)

General paper feedback

General small points

- No reason for an abstract or keywords
- Citations are good!
direct quotes need page number(s)

Abstract

Abstract

An abstract is a short summary of your (published or unpublished) research paper, usually about a paragraph (c. 6-7 sentences, 150-250 words) long. A well-written abstract serves multiple purposes:

- an abstract lets readers get the gist or essence of your paper or article quickly, in order to decide whether to read the full paper;
- an abstract prepares readers to follow the detailed information, analyses, and arguments in your full paper;
- and, later, an abstract helps readers remember key points from your paper.

From

<https://writing.wisc.edu/handbook/assignments/writing-an-abstract-for-your-research-paper/>

Abstract

Here are the typical kinds of information found in most abstracts:

1. the context or background information for your research; the general topic under study; the specific topic of your research
2. the central questions or statement of the problem your research addresses
3. what's already known about this question, what previous research has done or shown
4. the main reason(s), the exigency, the rationale, the goals for your research—Why is it important to address these questions? Are you, for example, examining a new topic? Why is that topic worth examining? Are you filling a gap in previous research? Applying new methods to take a fresh look at existing ideas or data? Resolving a dispute within the literature in your field? ...
5. your research and/or analytical methods
6. your main findings, results, or arguments
7. the significance or implications of your findings or arguments.

<https://writing.wisc.edu/handbook/assignments/writing-an-abstract-for-your-research-paper/>

Parts of a paper

- introduction
- method
- results
- discussion
- general discussion

Parts of a paper

- introduction
 - ▶ introduction to the overall topic
 - ▶ what is the research question(s) to be addressed
 - ▶ what will happen in the paper
- method
- results
- discussion
 - general discussion

Parts of a paper

- introduction
 - ▶ introduction to the overall topic
 - ▶ what is the research question(s) to be addressed
 - ▶ what will happen in the paper
- discussion
 - general discussion
- method
 - ▶ methods to approach the question(s) above
 - ▶ e.g. experiments
- results

Parts of a paper

- introduction
 - ▶ introduction to the overall topic
 - ▶ what is the research question(s) to be addressed
 - ▶ what will happen in the paper
- discussion
 - general discussion
- method
 - ▶ methods to approach the question(s) above
 - ▶ e.g. experiments
- results
 - ▶ what the did experiment(s) show

Parts of a paper

- introduction
 - ▶ introduction to the overall topic
 - ▶ what is the research question(s) to be addressed
 - ▶ what will happen in the paper
- method
 - ▶ methods to approach the question(s) above
 - ▶ e.g. experiments
- results
 - ▶ what the did experiment(s) show
- discussion
 - ▶ how does this help clarifying the research question(s)
 - ▶ relate results to the research question(s)
- general discussion

Parts of a paper

- introduction
 - ▶ introduction to the overall topic
 - ▶ what is the research question(s) to be addressed
 - ▶ what will happen in the paper
- method
 - ▶ methods to approach the question(s) above
 - ▶ e.g. experiments
- results
 - ▶ what the did experiment(s) show
- discussion
 - ▶ how does this help clarifying the research question(s)
 - ▶ relate results to the research question(s)
- general discussion
 - ▶ related the results and discussion to the overall topic
 - ▶ inverse of the introduction

Building models of statements

If the program returns the correct result, the code and compiler works.

The compiler did *not* work

Therefore the program did not return the correct result.

Building models of statements

If the program returns the correct result, the code and compiler works.

The compiler did *not* work

Therefore the program did not return the correct result.

$$A \rightarrow (B \wedge C)$$

$$\neg C$$

$$\therefore \neg A$$

What is an argument?

What is an argument?

What is an argument?

What is an argument?

an argument is a set of statements expressing the premises (whatever consists of empirical evidences and axiomatic truths) and an evidence-based conclusion.

What is an argument?

What is an argument?

an argument is a set of statements expressing the premises (whatever consists of empirical evidences and axiomatic truths) and an evidence-based conclusion.

When is an argument valid?

What is an argument?

What is an argument?

an argument is a set of statements expressing the premises (whatever consists of empirical evidences and axiomatic truths) and an evidence-based conclusion.

When is an argument valid?

An argument is valid if and only if it takes a form that make it impossible for the premises to be true and the conclusion nevertheless to be false.

if then statements: material conditional

What we say:

if p then q

What we write:

$p \rightarrow q$

if then statements: material conditional

What we say:

if p then q

p

What we write:

$p \rightarrow q$

p

if then statements: material conditional

What we say:

if p then q

p

therefore q

What we write:

$p \rightarrow q$

p

$\therefore q$

if then statements: material conditional

What we say:

if p then q

p

therefore q

What we write:

$p \rightarrow q$

p

$\therefore q$

Truth table

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

if then statements: material conditional

What we say:

if p then q

p

therefore q

What we write:

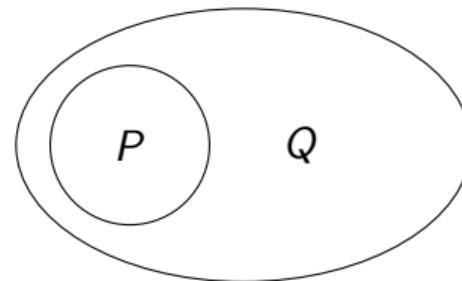
$p \rightarrow q$

p

$\therefore q$

Truth table

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T



if then statements: material conditional

$$p \rightarrow q$$

$$p$$

$$\therefore q$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

if then statements: material conditional

$$p \rightarrow q$$

$$p$$

$$\therefore q$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

if then statements: material conditional

$$\begin{array}{l} p \rightarrow q \\ p \\ \therefore q \end{array}$$

$$p \rightarrow q$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

if then statements: material conditional

$$\begin{aligned} p \rightarrow q \\ p \\ \therefore q \end{aligned}$$

$$\begin{aligned} p \rightarrow q \\ \neg p \end{aligned}$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

if then statements: material conditional

$$p \rightarrow q$$

$$p$$

$$\therefore q$$

$$p \rightarrow q$$

$$\neg p$$

?

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

if then statements: material conditional

$$p \rightarrow q$$

$$p$$

$$\therefore q$$

$$p \rightarrow q$$

$$\neg p$$

?

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

$$p \rightarrow q$$

if then statements: material conditional

$$p \rightarrow q$$

$$p$$

$$\therefore q$$

$$p \rightarrow q$$

$$\neg p$$

?

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- *if it rains then the street is wet*
- *if it is a Labrador then it is a dog*

$$\begin{array}{l} p \rightarrow q \\ \neg q \end{array}$$

if then statements: material conditional

$$\begin{aligned} p \rightarrow q \\ p \\ \therefore q \end{aligned}$$

$$\begin{aligned} p \rightarrow q \\ \neg p \\ ? \end{aligned}$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- if it rains *then* the street is wet
- if it is a Labrador *then* it is a dog

$$\begin{aligned} p \rightarrow q \\ \neg q \\ \therefore \neg p \end{aligned}$$

if then statements: material conditional

$$\begin{aligned} p \rightarrow q \\ p \\ \therefore q \end{aligned}$$

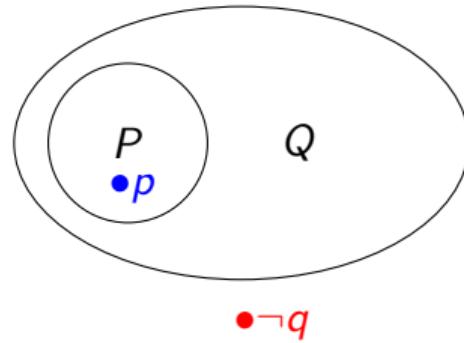
$$\begin{aligned} p \rightarrow q \\ \neg p \\ ? \end{aligned}$$

p	q	$p \rightarrow q$
T	T	T
T	F	F
F	T	T
F	F	T

Examples:

- if it rains *then* the street is wet
- if it is a Labrador *then* it is a dog

$$\begin{aligned} p \rightarrow q \\ \neg q \\ \therefore \neg p \end{aligned}$$



Truth tables

\wedge (and)

p	q	$p \wedge q$
T	T	T
T	F	F
F	T	F
F	F	F

Truth tables

\wedge (and)

p	q	$p \wedge q$
T	T	T
T	F	F
F	T	F
F	F	F

\vee (or)

p	q	$p \vee q$
T	T	T
T	F	T
F	T	T
F	F	F

Truth tables

\wedge (and)

p	q	$p \wedge q$
T	T	T
T	F	F
F	T	F
F	F	F

\vee (or)

p	q	$p \vee q$
T	T	T
T	F	T
F	T	T
F	F	F

\oplus (Exclusive-or)

p	q	$p \oplus q$
T	T	F
T	F	T
F	T	T
F	F	F

Truth tables

\wedge (and)

p	q	$p \wedge q$
T	T	T
T	F	F
F	T	F
F	F	F

\vee (or)

p	q	$p \vee q$
T	T	T
T	F	T
F	T	T
F	F	F

\oplus (Exclusive-or)

p	q	$p \oplus q$
T	T	F
T	F	T
F	T	T
F	F	F

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

Example

If the program returns the correct result, the code and compiler works.

The compiler did *not* work

Therefore the program did not return the correct result.

Example

If the program returns the correct result, then
the code and compiler works.

The compiler did not work

Therefore the program did not return the
correct result.

Example

If the program returns the correct result, then
the code and compiler works.

The compiler did not work

Therefore the program did not return the
correct result.

$$A \rightarrow (B \wedge C)$$

$$\neg C$$

$$\therefore \neg A$$

Example

If the program returns the correct result, then
the code and compiler works.

The compiler did not work

Therefore the program did not return the
correct result.

$$A \rightarrow (B \wedge C)$$

$$\neg C$$

$$\therefore \neg A$$

		\wedge (and)	
B	C	$(B \wedge C)$	
T	T	T	
T	F	F	
F	T	F	
F	F	F	

A	$(B \wedge C)$	$A \rightarrow (B \wedge C)$
T	T	T
T	F	F
F	T	T
F	F	T

Questions?

1. Machine learning
2. Encoding
 - example: rapid tuning shifts in human auditory cortex ...
3. Decoding
 - evaluating algorithms
 - linear vs non-linear models
4. Time generalisation
 - example: brain mechanisms underlying the brief maintenance ...
5. General paper feedback
 - abstract
 - parts of a paper
 - arguments
 - material condition

References I

- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.
<https://doi.org/10.1023/A:1010933404324>
- Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD '16*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Hastie, T., Tibshirani, R., Friedman, J., Hastie, T., Friedman, J., & Tibshirani, R. (2009). *The elements of statistical learning* (Vol. 2). Springer.
<http://statweb.stanford.edu/~tibs/ElemStatLearn/>
- Holdgraf, C. R., de Heer, W., Pasley, B., Rieger, J., Crone, N., Lin, J. J., Knight, R. T., & Theunissen, F. E. (2016). Rapid tuning shifts in human auditory cortex enhance speech intelligibility. *Nat Commun*, 7, 13654. <https://doi.org/10.1038/ncomms13654>
- Holdgraf, C. R., Rieger, J. W., Micheli, C., Martin, S., Knight, R. T., & Theunissen, F. E. (2017). Encoding and Decoding Models in Cognitive Electrophysiology. *Front Syst Neurosci*, 11, 61. <https://doi.org/10.3389/fnsys.2017.00061>

References II

- King, J. R., & Dehaene, S. (2014). Characterizing the dynamics of mental representations: The temporal generalization method. *Trends Cogn Sci*, 18(4), 203–10.
<https://doi.org/10.1016/j.tics.2014.01.002>
- King, J. R., Gwilliams, L., Holdgraf, C., Sassenhagen, J., Barachant, A., Engemann, D., Larson, E., & Gramfort, A. (2018). Encoding and Decoding Neuronal Dynamics: Methodological Framework to Uncover the Algorithms of Cognition. 19.
<https://hal.archives-ouvertes.fr/hal-01848442>
- King, J. R., Pescetelli, N., & Dehaene, S. (2016). Brain Mechanisms Underlying the Brief Maintenance of Seen and Unseen Sensory Information. *Neuron*, 92(5), 1122–1134.
<https://doi.org/10.1016/j.neuron.2016.10.051>
- Raizada, R. D. S., & Lee, Y.-S. (2013). Smoothness without Smoothing: Why Gaussian Naive Bayes Is Not Naive for Multi-Subject Searchlight Studies. *PLoS ONE*, 8, e69566.
<https://doi.org/10.1371/journal.pone.0069566>

References III

- Varoquaux, G., Raamana, P. R., Engemann, D. A., Hoyos-Idrobo, A., Schwartz, Y., & Thirion, B. (2017). Assessing and tuning brain decoders: Cross-validation, caveats, and guidelines. *NeuroImage*, 145, 166–179.
<https://doi.org/10.1016/j.neuroimage.2016.10.038>
- Westner, B. U., Dalal, S. S., Hanslmayr, S., & Staudigl, T. (2018). Across-subjects classification of stimulus modality from human MEG high frequency activity (D. Bush, Ed.). *PLOS Computational Biology*, 14(3), e1005938.
<https://doi.org/10.1371/journal.pcbi.1005938>