

Artificial neural networks

Mads Jensen, PhD

✉ mads@cas.au.dk

November 16th, 2020



AARHUS UNIVERSITY



INTERACTING MINDS CENTRE



Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

1 / 41

Contents

1. Artificial neural networks

- Linear layer
- Activation functions
- Loss/cost function
- Learning

2. Deep learning

3. Convolutional neural networks

- Convolution
- Max pooling

4. Adversarial attacks

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

2 / 41

Artificial neural networks

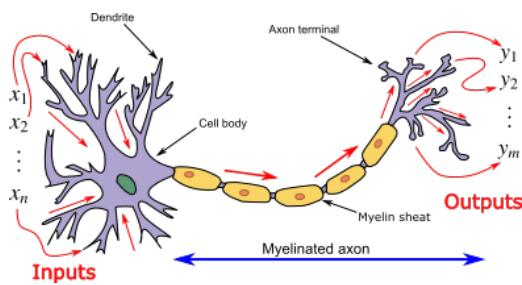
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

3 / 41

(Artificial) neurons



Real neuron:

- connected to other neurons
- can perform simple computations
- inhibition and exhibition

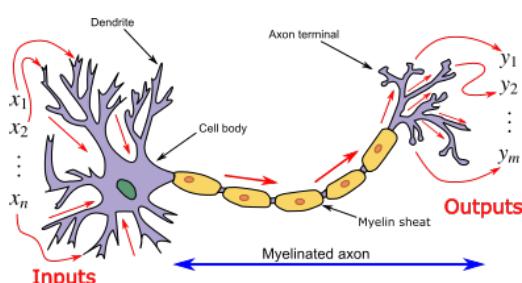
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

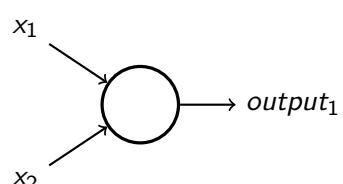
4 / 41

(Artificial) neurons



Real neuron:

- connected to other neurons
- can perform simple computations
- inhibition and exhibition



Artificial neuron:

- connected to other neurons
- can perform simple computations
- inhibition and exhibition

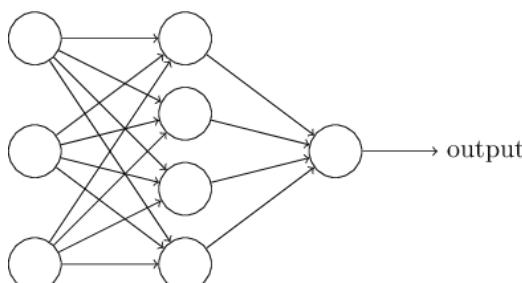
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

4 / 41

Neural network



(Figure from Nielsen, 2015)

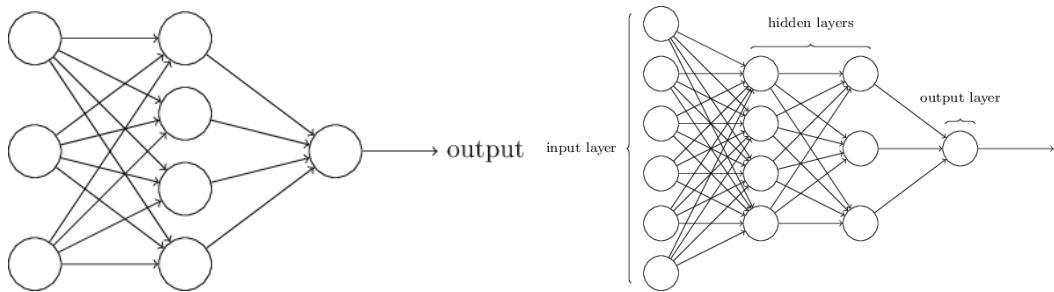
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

5 / 41

Neural network



(Figure from Nielsen, 2015)

XOR (exclusive or)

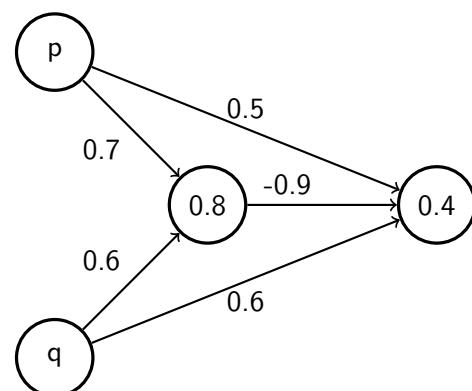
p	q		$p \oplus q$
T	T		F
T	F		T
F	T		T
F	F		F

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$

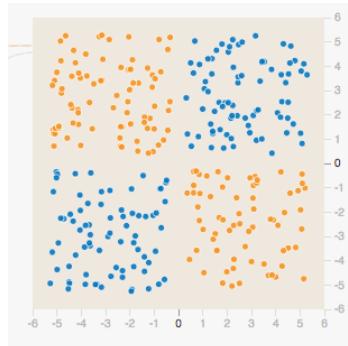
XOR (exclusive or)

p	q		$p \oplus q$
T	T		F
T	F		T
F	T		T
F	F		F

$$p \oplus q = (p \vee q) \wedge \neg(p \wedge q)$$



XOR (exclusive or)



(Figure from <https://playground.tensorflow.org/>)

Machine learning

Supervised learning

Data $\{X, y\}$

Model $y \approx f_\theta(X)$

$$\text{Loss/cost} \quad \quad \mathcal{L}(\theta) = \sum_{i=1}^N l(f_\theta(x_i), y_i)$$

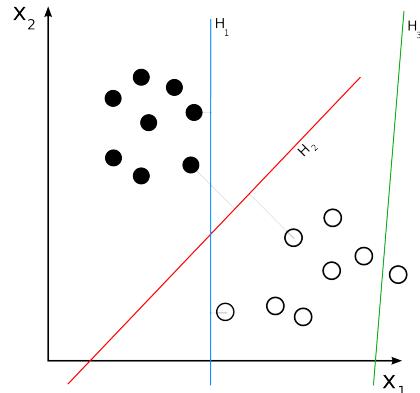
$$\text{Optimisation} \quad \theta^* = \arg \min_{\theta} \mathcal{L}(\theta)$$

Neural network parts

- Linear layer
 - Activation function
 - Loss/cost function
 - Neurons: units
 - Parameters: weights

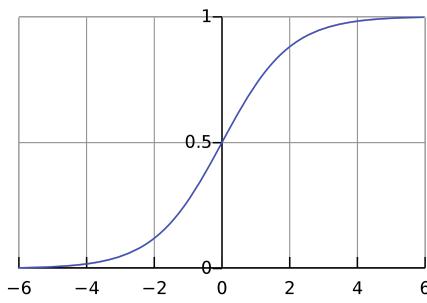
Linear layer

$$f_{\text{linear}}(x, W, b) = Wx + b$$



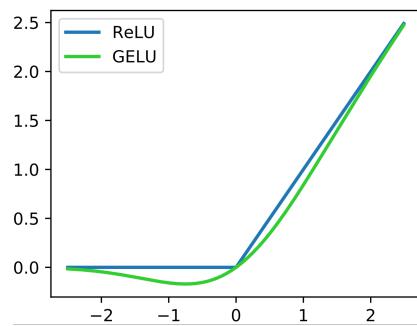
(Figure from https://en.wikipedia.org/wiki/Linear_classifier)

Activation functions



Sigmoid function

$$S(x) = \frac{1}{1+e^{-x}}$$



Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$

(Figure from [https://en.wikipedia.org/wiki/Rectifier_\(neural_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks)))

Softmax

Softmax is a generalization of the logistic function to multiple dimensions.

$$f_{sm}(x) = \frac{e^x}{\sum_{j=1}^k e^{x_j}}$$

Often used as the last layer before the loss/cost function

Loss/cost function

Cross entropy (also called logarithmic loss, log loss, or logistic loss)

$$H(p, q) = - \sum_i p_i \log_2(q_i)$$

Loss/cost function

Cross entropy (also called logarithmic loss, log loss, or logistic loss)

$$H(p, q) = - \sum_i p_i \log_2(q_i)$$

Shannon entropy (from lecture 5):

$$H(X) = - \sum_{i=1}^n p(x_i) \log_2 p(x_i)$$

Learning

- No learning yet!

Learning

- No learning yet!
- Building block for learning
 - ▶ Function ($y \approx f_\theta(X)$)
 - ▶ Loss/cost
 - ▶ Optimisation
 - ▶ Data

Learning

- No learning yet!
- Building block for learning
 - ▶ Function ($y \approx f_\theta(X)$)
 - ▶ Loss/cost
 - ▶ Optimisation
 - ▶ Data

What is learning in this context?

Learning

- No learning yet!
- Building block for learning
 - ▶ Function ($y \approx f_\theta(X)$)
 - ▶ Loss/cost
 - ▶ Optimisation
 - ▶ Data

What is learning in this context?

- Changing response/output to a fixed stimulus

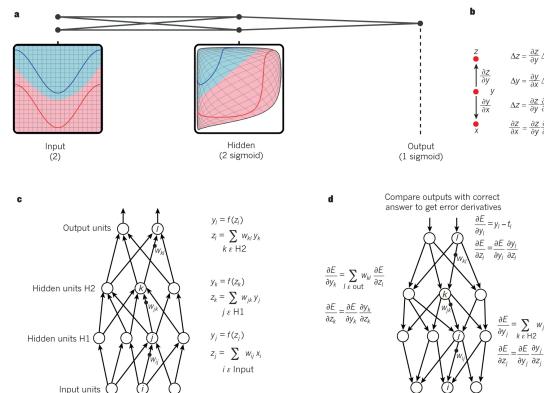
Learning

- No learning yet!
- Building block for learning
 - ▶ Function ($y \approx f_\theta(X)$)
 - ▶ Loss/cost
 - ▶ Optimisation
 - ▶ Data

What is learning in this context?

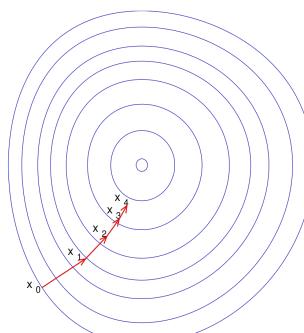
- Changing response/output to a fixed stimulus
- Response *after* seeing data is not the same as *before* seeing data.

Backpropagation



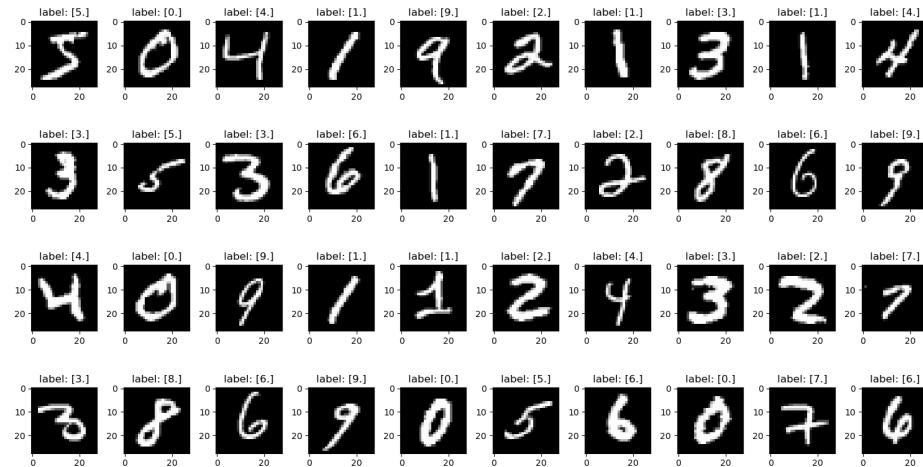
(Figure from LeCun et al., 2015)

Gradient descent



(Figure from https://en.wikipedia.org/wiki/Gradient_descent)

Mini batches



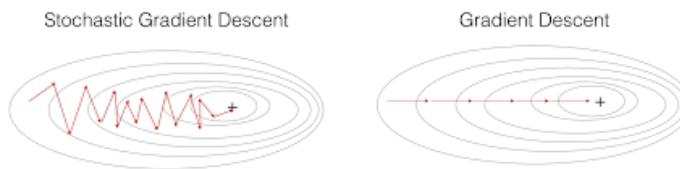
Data from <http://yann.lecun.com/exdb/mnist/>

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 17 / 41

Stochastic gradient decent



(Figure from <https://medium.com/bayshore-intelligence-solutions/why-is-stochastic-gradient-descent-2c17baf016de>)

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 18 / 41

Deep learning

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 19 / 41

Deep learning

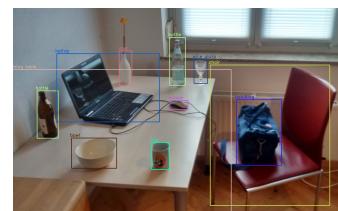
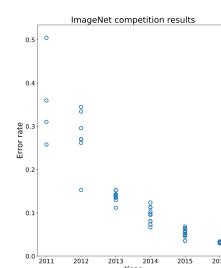
- Deep learning
- Deep neural networks

Deep learning

- Deep learning
- Deep neural networks
- Deep vs shallow neural networks
 - ▶ Shallow ≤ 3 hidden layers
 - ▶ Deep > 3 hidden layers

Deep learning

- Deep learning
- Deep neural networks
- Deep vs shallow neural networks
 - ▶ Shallow ≤ 3 hidden layers
 - ▶ Deep > 3 hidden layers



(Figures from https://en.wikipedia.org/wiki/Object_detection and <https://en.wikipedia.org/wiki/ImageNet>)

Technologies



Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 21 / 41

Technologies



(Figure from <https://upload.wikimedia.org/wikipedia/en/a/a7/Quake-gameplay.png>)

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 21 / 41

Technologies

- Graphical processing unit (GPU)



Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 22 / 41

Technologies

- Graphical processing unit (GPU)
- Data
 - ▶ Internet
 - ▶ Facebook, Twitter etc.
 - ▶ Deep learning need a lot of training data

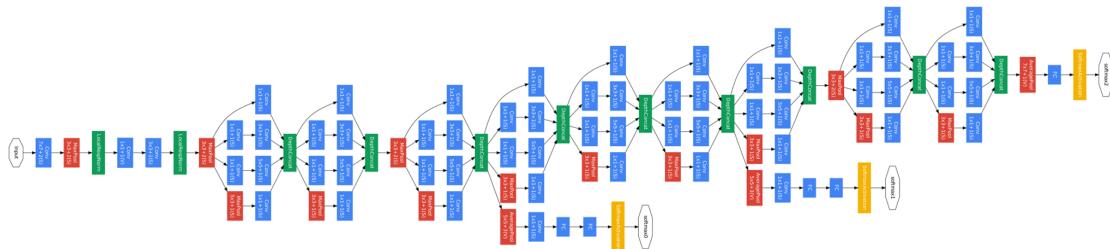


Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 22 / 41

GoogLeNet



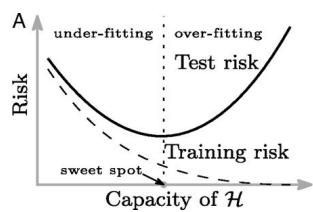
(Figure from Szegedy et al., 2015)

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 23 / 41

Underfitting & overfitting



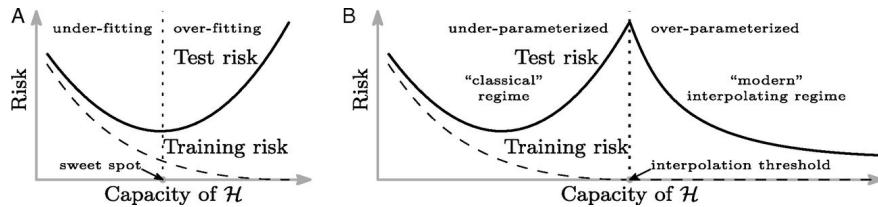
(Figure from Belkin et al., 2019)

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 24 / 41

Underfitting & overfitting



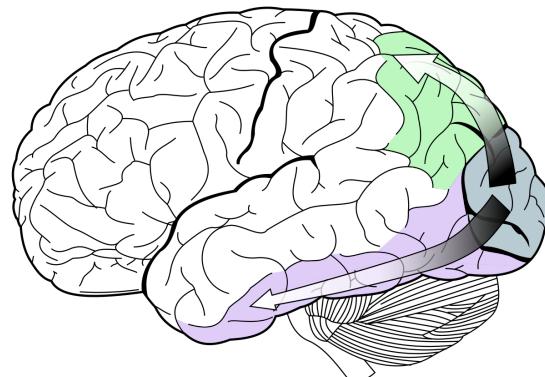
(Figure from Belkin et al., 2019)

Convolutional neural networks



(Figure from <https://en.wikipedia.org/wiki/Tree>)

Convolutional neural networks



(Figure from https://en.wikipedia.org/wiki/Visual_perception)

Mads Jensen (RFR, IMC, & CFIN)

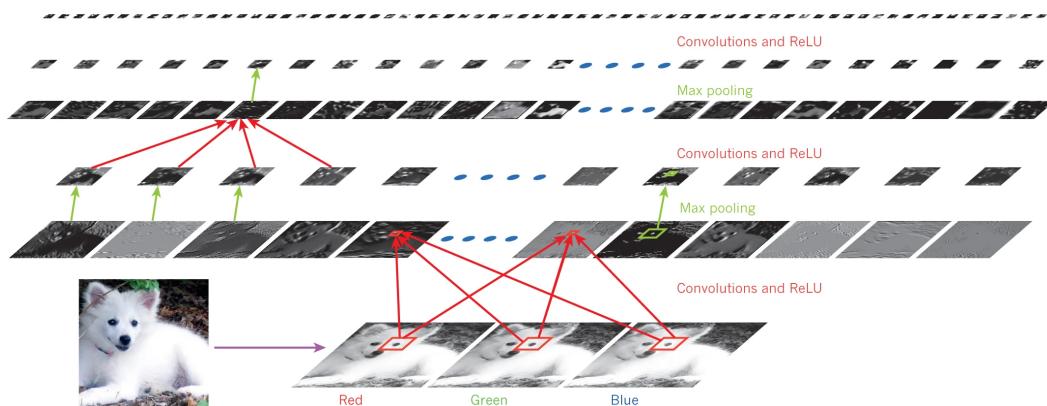
Artificial neural networks

November 16th, 2020

27 / 41

Convolutional neural networks

Samoyed (16); Papillon (5.7); Pomeranian (2.7); Arctic fox (1.0); Eskimo dog (0.6); white wolf (0.4); Siberian husky (0.4)



(Figure from LeCun et al., 2015)

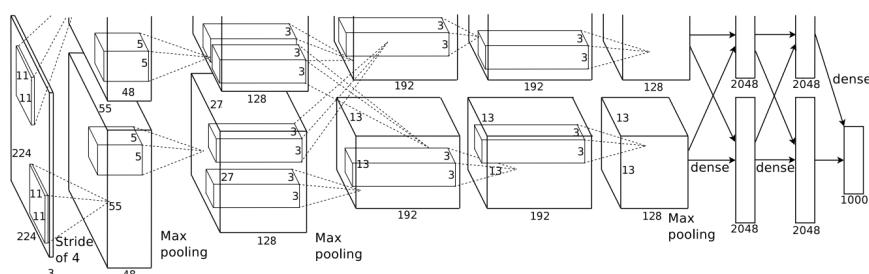
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

28 / 41

AlexNet



(Figure from Krizhevsky et al., 2017)

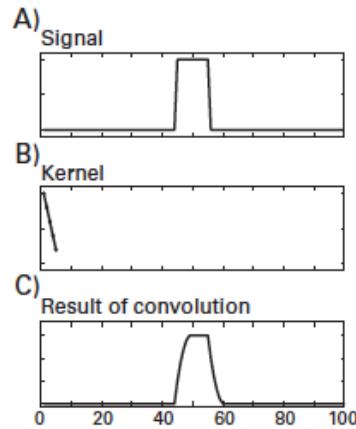
Mads Jønson (RER, IMC & CEIN)

Artificial neural networks

November 16th, 2020

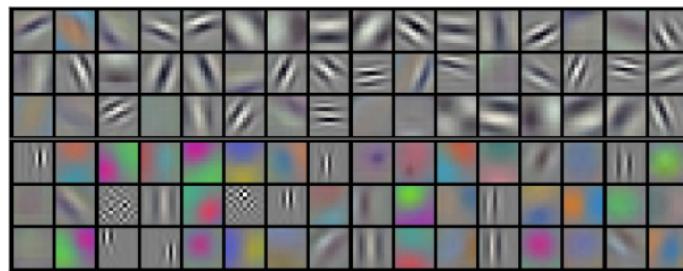
39 / 41

Convolution



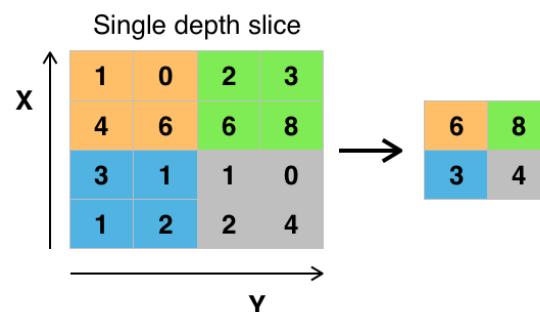
(Figure from Cohen, 2014)

Convolution kernels



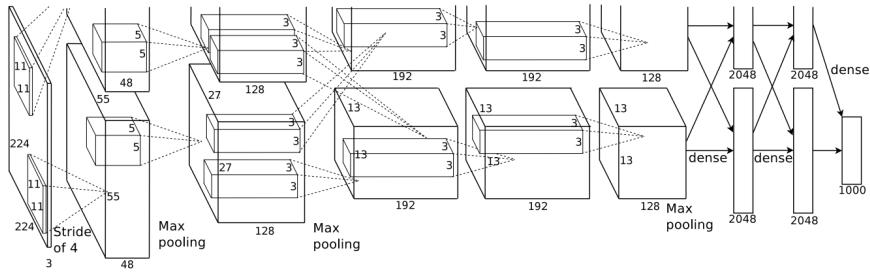
(Figure from Krizhevsky et al., 2017)

Max pooling



(Figure from https://en.wikipedia.org/wiki/Convolutional_neural_network)

AlexNet



(Figure from Krizhevsky et al., 2017)

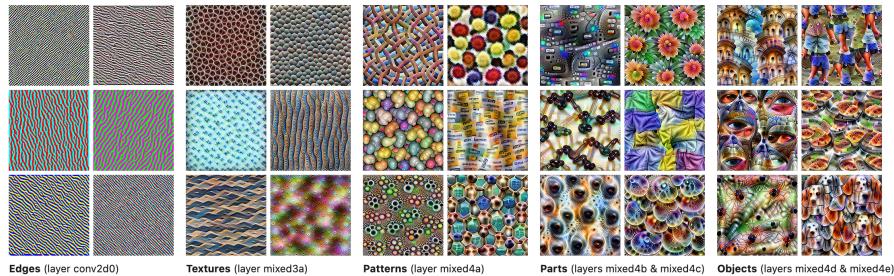
Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

33 / 41

Feature visualisation



(Figure from Olah et al., 2017)

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020 34 / 41

Problems with max pooling

The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster.

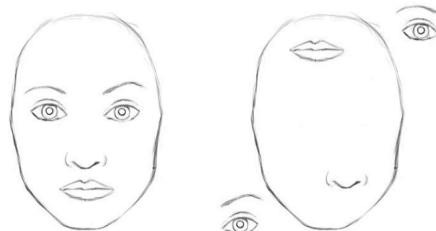
Hinton, reddit AMA¹

https://www.reddit.com/r/MachineLearning/comments/2lmo0l/ama_geoffrey_hinton/

Problems with max pooling

The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster.

Hinton, reddit AMA¹

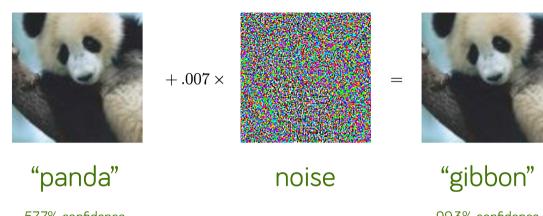


(Figure from <https://medium.com/ai%C2%B3-theory-practice-business/understanding-hintons-capsule-networks-part-i-intuition-b4b559d1159b>)

¹https://www.reddit.com/r/MachineLearning/comments/2lmo0l/ama_geoffrey_hinton/ Mads Jensen (RFR, IMC, & CFIN) Artificial neural networks November 16th, 2020 35 / 41

Adversarial attacks

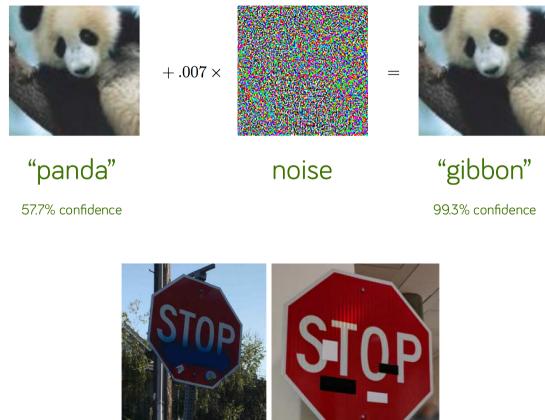
Adversarial attacks



(Top figure from Goodfellow et al (2015);

Mads Jensen (RFR, IMC, & CFIN) Artificial neural networks November 16th, 2020 37 / 41

Adversarial attacks



(Top figure from Goodfellow et al (2015); bottom from Eykholt et al. (2018))

A set of small, light-gray navigation icons typically found in presentation software like Beamer. They include symbols for back, forward, search, and other document-related functions.

“Assignment” for next lecture

What to do in last lecture?

In the syllabus for the last lecture:

- summary
 - Q & A

Consider if there is anything you would like to have covered

A set of small, light-gray navigation icons typically found in presentation software like Beamer. From left to right, they include: a left arrow, a square, a right arrow, a double left arrow, a double square, a double right arrow, a double left arrow with a horizontal line, a double right arrow with a horizontal line, a vertical ellipsis, a circular arrow, a magnifying glass, and a circular arrow with a dot.

Questions?

1 Artificial neural networks

- Linear layer
 - Activation functions
 - Loss/cost function
 - Learning

2. Deep learning

3. Convolutional neural networks

- Convolution
 - Max pooling

4. Adversarial attacks

Mads Jensen (RFR, IMC, & CFIN)

Artificial neural networks

November 16th, 2020

39 / 41

References I

- Belkin, M., Hsu, D., Ma, S., & Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32), 15849–15854. <https://doi.org/10.1073/pnas.1903070116>
- Cohen, M. X. (2014). *Analyzing neural time series data: Theory and practice*. The MIT Press.
- Eykholt, K., Evtimov, I., Fernandes, E., Li, B., Rahmati, A., Xiao, C., Prakash, A., Kohno, T., & Song, D. (2018). Robust Physical-World Attacks on Deep Learning Visual Classification. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 1625–1634. <https://doi.org/10.1109/CVPR.2018.00175>
- Goodfellow, I. J., Shlens, J., & Szegedy, C. (2015, March 20). *Explaining and Harnessing Adversarial Examples*. arXiv: 1412.6572 [cs, stat]. Retrieved November 15, 2020, from <http://arxiv.org/abs/1412.6572>
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2017). ImageNet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6), 84–90. <https://doi.org/10.1145/3065386>

References II

- LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521, 436–444. <https://doi.org/10.1038/nature14539>
- Nielsen, M. A. (2015). *Neural networks and deep learning* (Vol. 2018). Determination press San Francisco, CA. <http://neuralnetworksanddeeplearning.com/>
- Olah, C., Mordvintsev, A., & Schubert, L. (2017). Feature visualization. *Distill*. <https://doi.org/10.23915/distill.00007>
- Szegedy, C., Wei Liu, Yangqing Jia, Sermanet, P., Reed, S., Anguelov, D., Erhan, D., Vanhoucke, V., & Rabinovich, A. (2015). Going deeper with convolutions. *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1–9. <https://doi.org/10.1109/CVPR.2015.7298594>